Predicting Loan Payback

Ziyan Chen · October 2026 · Repository: github.com/ZiyanChen1/predicting-loan-payback

Contents. (1) Topic exploration · (2) One-page topic and industry infographic · (3) Selected focus and project proposal · (4) Initial data exploration · (5) Baseline and proof of concept · (6) Open questions and next steps · Key terms · References

1. Topic exploration: what each topic could do for a real organization

I considered all six topic areas. For each, one realistic example and the decision it could support:

Topic area	Realistic example	What the organization could learn or do
Supervised machine learning	**Banking.** A bank predicts whether a borrower will repay a loan, using income, debt-to-income ratio, credit score, and loan terms.	Rank borrowers by risk, set interest rates, and send borderline files to a loan officer.
Unsupervised machine learning	**Retail.** A supermarket groups loyalty-card customers by what, when, and how often they buy.	Give each group a different discount to increase sales, for example bulk coupons for family shoppers.
Dimensionality reduction	**Manufacturing.** A factory uses PCA to combine 40 related machine sensor readings into 2–3 summary scores.	Spot machines drifting away from normal and repair them before they break. Kept: the overall pattern. Lost: which single sensor caused it.
Natural language processing	**Telecom.** A phone company sorts customer support messages into billing, network outage, or cancellation.	Send each message to the right team and track which problems grow each week.
Semantic or sentiment analysis	**Hospitality.** A hotel chain scores guest reviews as positive or negative for room, staff, and breakfast.	Find the service that upsets guests most and fix it first.
Social media / digital channel analytics	**Consumer brand.** A coffee chain compares the likes, shares, and reach of its own Instagram posts.	Choose which posts (new drinks, promotions) and posting times to use next.

Credit scoring: from expert judgment to explainable machine learning

Supervised classification has been used to predict loan repayment for nearly 70 years. What changed is the model, and what never changed is the duty to explain.

Topic: supervised ML
Industry: consumer lending

ORIGINS AND DEVELOPMENT

Scoring industry · Machine learning · Regulation

1956
Fair, Isaac and Company, known today as FICO, is founded [11].

Late 1950s
Credit scoring first emerges to support credit decisions at large retail stores and finance companies; by the end of the 1970s, most of the largest US commercial banks use scoring systems [1].

1974
The Equal Credit Opportunity Act (ECOA). Under Regulation B, reasons given for an adverse action "must be specific and indicate the principal reason(s)" [2].

1989
FICO and the three credit bureaus release the FICO Score. Generic, bureau-based scores that work across many loan products appear at the end of the 1980s [1] [11].

2004
Basel II lets banks use their own internal rating models to estimate the probability of default for capital rules [3].

2016–2017
XGBoost and LightGBM make gradient-boosted trees fast and accurate on large tables [4][5]. SHAP explains each prediction feature by feature [6].

2021
FinRegLab: machine learning underwriting is "intensifying" in US credit markets, most advanced in credit cards and unsecured personal loans [7].

2022
CFPB Circular 2022-03: using a complex algorithm does not excuse a lender from giving specific reasons [8]. (Withdrawn in May 2025; the Regulation B requirement itself remains [2][8].)

2024
U.S. Treasury: AI "has long been deployed" in credit underwriting; black-box models and possible bias remain key concerns [12].

PRACTICAL USE TODAY

52% of surveyed banks use ML in production for credit risk [9]

25% more are piloting ML for credit risk [9]

43 banks in 9 regions surveyed, including the U.S. (2022) [9]

WHO USES IT, AND FOR WHAT

- **Lenders using credit scores:** to approve applications, set interest rates by risk, manage accounts, and estimate losses [1].
- **Banks:** scorecards for approval; internal rating models to estimate default risk for capital [3].
- **Card and personal-loan lenders, fintechs:** ML underwriting that evaluates tens of thousands of applicants each week [7].
- Explainability is a primary challenge for ML in credit risk [9], so ML is often tested as a **challenger** next to a simpler scorecard.

VALUE AND IMPACT

POSSIBLE BENEFITS

Better risk ranking can mean fewer defaults at the same loan volume, or more approvals at the same loss; faster and more consistent decisions [7].

COSTS AND RISKS

Lenders must still give specific reasons [2]; black-box models give limited insight, and poorly trained AI can reinforce bias [12]; complex models are harder to validate.

Benefits depend on the lender and data; they are not guaranteed savings.

CONNECTION TO MY PROJECT

- **Champion vs. challenger:** the data includes the lender's own grade (A1 to F5), so I can test ML against the lender's current ranking, like a real model review.
- **Method:** logistic regression (transparent) and LightGBM (gradient boosting [5]).
- **Explainable:** feature contributions (the same idea as SHAP [6]) give each borrower's main risk drivers.
- **Early evidence:** AUC 0.922 vs. 0.655 for the lender's grade (synthetic data; see Section 5).

[1] Federal Reserve Board (2007), Report to the Congress on credit scoring.
[2] CFPB, Regulation B, 12 CFR §1002.9.
[3] Basel Committee on Banking Supervision (2004), Basel II revised framework.
[4] Chen & Guestrin (2016), XGBoost, KDD.
[5] Ke et al. (2017), LightGBM, NeurIPS.
[6] Lundberg & Lee (2017), SHAP, NeurIPS.
[7] FinRegLab (2021), ML for credit underwriting: market context.
[8] CFPB (2022), Circular 2022-03; withdrawal: Morgan Lewis CFPB guidance tracker.
[9] IIF & EY (2022), Survey on machine learning in credit risk.
[11] myFICO (2018), The history of the FICO Score.
[12] U.S. Treasury (Dec. 2024), AI in financial services report.
Full links in the reference list at the end of this document.

3. Selected focus and project proposal

> **Question.** Can machine learning rank approved borrowers by default risk better than the lender's own grade, by enough to matter, while staying explainable?

Element	My project
Problem	About 1 in 5 approved loans is not repaid. The lender ranks borrowers with its own grade (A1 to F5). A better ranking leaves less money in loans that default.
Stakeholder	A consumer lender's credit risk manager (sets policy) and loan officers (review borderline files).
Label	`loan_paid_back`: 1 = repaid, 0 = not repaid. Binary classification. Positive class for errors = **not repaid**.
Action	Route similar future applicants: **approve / review / reject** (low / medium / high risk); evaluated by operational cost.
Data	Kaggle Playground Series S5E11, "Predicting Loan Payback" [10].
Method	Logistic regression and gradient boosting (LightGBM), compared with three baselines.

Intended use

- **How much better:** defaulted principal left if the lender kept only its safest 70 / 80 / 90% of loans.
- **What drives the risk:** top three risk drivers for each high-risk borrower (model explanations, not decline reasons).
- **How to use it:** approve / review / reject cutoffs chosen with business costs, plus a sensitivity check on the cost assumptions.

Success measures

Measure	Useful if
ROC AUC (ranking)	Above the lender's grade (0.655) and a classic credit score + DTI scorecard (0.768).
Recall on loans not repaid	Of the loans that default, how many are caught? Higher than every baseline.
Operational cost	Less defaulted principal than the lender's grade at the same share of loans kept.
Interpretability	Each high-risk borrower gets readable risk drivers that agree with the EDA.

Dataset and access

- **Access:** sign in to Kaggle, accept the competition rules, and download the files into `data/` (steps in the README). The rules do not allow redistributing the data, so the public repository holds code, figures, and results only.
- **One row** = one loan that was already approved: the borrower's finances, loan terms, profile, and whether it was paid back.

Key field	Business meaning	Expected signal
employment_status	Stability of income	Unemployed or student → much higher risk
debt_to_income_ratio	Existing debt payments ÷ income	Higher DTI → higher risk
credit_score	Bureau score (395 to 849)	Higher score → lower risk
grade_subgrade	Lender's own grade, A1 (best) to F5 (worst)	The lender's current ranking
annual_income, loan_amount, interest_rate	Ability to repay, exposure, price of risk	Numeric features
marital_status, education_level, loan_purpose	Profile and use of funds	Weak signals
gender	Borrower gender	Excluded; fairness check only

Limits. (1) **Synthetic:** generated from an original loan dataset, so dollars show direction, not a real bank's results. (2) **Approved loans only:** declined applicants are not observed. (3) **Privacy:** no personal identifiers. (4) **Licensing:** Kaggle rules prohibit sharing the data with non-participants.

Planned method

Step	What I do
Shortlist	For tabular data: a logistic model plus a tree ensemble.
Baselines	(1) Majority baseline: always predict "repaid" (high accuracy, zero recall); (2) the lender's grade only; (3) a classic credit score + DTI scorecard.
Logistic regression	A transparent first challenger: $z = b + w \cdot x$, $p = 1 / (1 + e^{-z})$, trained with log loss (maximum likelihood). Coefficients show the direction and size of each risk factor.
Gradient boosting	LightGBM adds trees sequentially; a strong tabular model that can learn the nonlinear DTI threshold.
Preprocessing	Learn the preparation once from training rows and reuse it: ColumnTransformer → Pipeline → cross-validation. The scaler and encoder are fit only on the training partition to prevent leakage.
Evaluation	Fit on training data, choose on validation, report on untouched test data: stratified 5-fold CV to compare models, a fixed 20% validation set to choose cutoffs, and the Kaggle test set scored once.

4. Initial data exploration

Check	Result	Decision
Size	593,994 training rows × 13 columns; 254,569 test rows	Enough for a fixed 20% validation set
Missing / duplicates	None	No imputation needed
Target balance	79.9% repaid, 20.1% not repaid	Accuracy is misleading; use ROC AUC and recall
Categories	No spelling or case clashes; same levels in train and test	Keep "Other" as a real category
Outliers (Tukey 1.5 × IQR)	About 3% of income and DTI values	Keep: high-DTI outliers repaid 45% vs. 81%; they are real high-risk borrowers
Skewness	Income and DTI right-skewed	Log transform
Correlation	Credit score vs. grade: r = −0.91	Logistic regression keeps credit score only

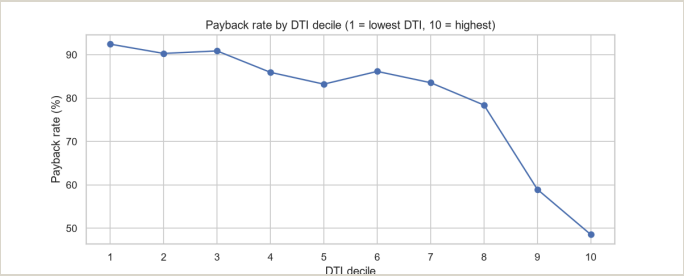


Figure 1. Payback rate by DTI decile.

Interpretation. Payback is 83–92% in the first seven deciles, then falls to 78%, 59%, and 49%. Risk is nonlinear, so a tree model is a good fit.

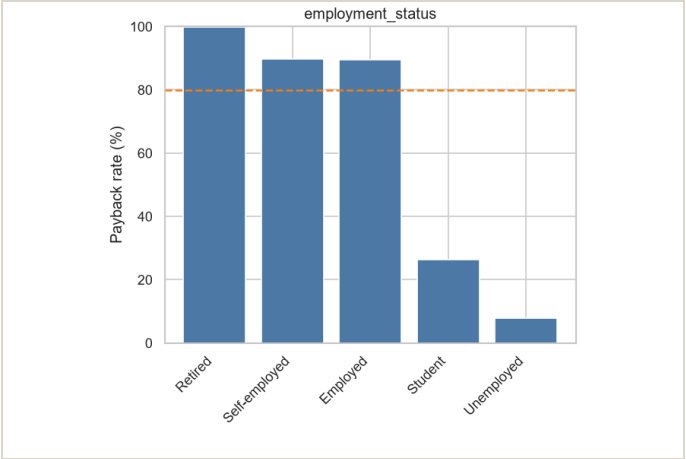


Figure 2. Payback rate by employment status (dashed line = 79.9%).

Interpretation. The strongest signal: 7.8% of unemployed borrowers repaid vs. 89.4% of employed. Other categories barely change the payback rate.

Full audit: `01_eda.ipynb`, `src/eda_audit.py`, and `reports/`.

5. Baseline and proof of concept

Stratified 5-fold cross-validation on the training data (`02_modeling.ipynb`):

Model	ROC AUC	Recall (not repaid)
Majority baseline: always "repaid"	0.500	0.0%
Lender's grade only	0.655	1.8%
Classic scorecard: credit score + DTI	0.768	19.3%
Logistic regression	0.908	59.8%
LightGBM	**0.922**	**60.9%**

Recall at the default 0.5 threshold. The majority baseline is 79.9% accurate but catches no defaults.

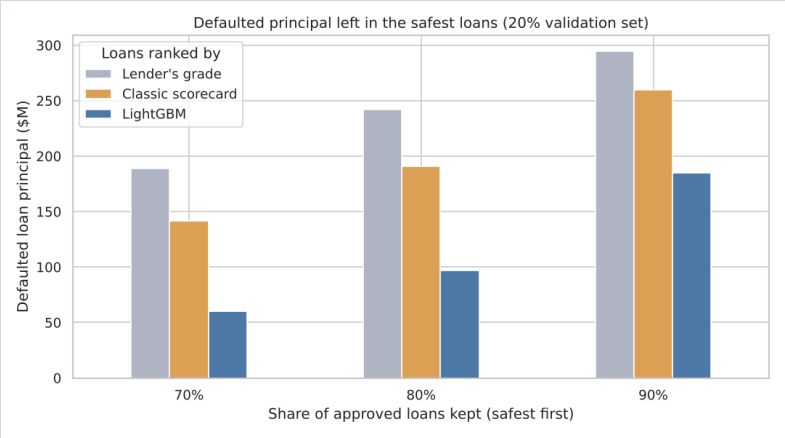


Figure 3. Defaulted principal left in the safest loans (20% validation set, 118,799 loans).

Interpretation. Keeping the safest 80% of approved loans, LightGBM leaves **$97.1M** of defaulted principal vs. **$242.0M** with the lender's grade: about **60% less**. The gap holds at 70% and 90%.

Synthetic data: the direction matters more than the exact dollars.

First result	Illustrative
Risk drivers	Riskiest 10% of borrowers (97.6% defaulted): top drivers are unemployed, high DTI, low credit score.
Approve / review / reject	Assumed costs: 60% loss given default, $50 per review, review capacity 15%. Cutoffs: approve ≥ 0.89, review 0.77–0.89, reject < 0.77. Default rates: 3.7% / 15.4% / 67.1%.

Feasibility. The models run in minutes and already meet the AUC and dollar success measures.

6. Open questions and next steps

Still unknown

- **Calibration:** a strong ROC AUC does not prove that a predicted 80% means 80% in reality.
- **Costs:** loss given default, review cost, and capacity are assumptions.
- **Fairness:** do errors differ across gender groups?
- **Real data:** does a model trained on synthetic rows hold up on the original loan dataset?

Next steps

Step	What I will do
Confusion matrix	Name each kind of mistake at the chosen cutoffs; report precision and recall with counts.
Thresholds	Finalize approve / review / reject cutoffs on the validation set with business costs.
Calibration and fairness	Calibration curve; compare errors across gender groups.
Tuning	Tune LightGBM with cross-validation only.
Model handoff	Save the fitted pipeline (joblib) with its preprocessing and metadata; reload it and reproduce a reference prediction.
Local API	FastAPI `/predict`: one applicant in; payback probability, risk tier, and top drivers out; invalid input returns 422; packaged in Docker.
Communicate	Data-to-decision pipeline visual, written report, and stakeholder presentation.

Key terms

ROC AUC: How well the model ranks borrowers who repay above those who default. 0.5 = random guess; 1.0 = perfect.

Accuracy: Share of all predictions that are correct. Misleading when one outcome dominates (80% of loans are repaid).

Recall (not repaid): Of the loans that actually defaulted, the share the model flagged.

Majority baseline: A model that always predicts the most common outcome ("repaid"). Any useful model must beat it.

Lender's grade: The lender's own risk rating, A1 (safest) to F5 (riskiest); the current way of ranking borrowers.

Scorecard: A simple, transparent credit model built from a few factors, such as credit score and DTI.

DTI: Debt-to-income ratio: existing debt payments divided by income.

Defaulted principal: Loan amount that was not repaid, before any recovery.

Loss given default: Share of a defaulted loan the lender actually loses after recovery.

Logistic regression: A model that turns a weighted sum of factors into a probability between 0 and 1.

Gradient boosting / LightGBM: Many small decision trees added one after another, each correcting earlier errors. LightGBM is a fast version.

Cross-validation (5-fold): Split the training data into 5 parts; train on 4, test on 1, and rotate, so every row is tested once.

Validation set: Data held back from training and used to choose settings such as cutoffs.

Data leakage: Information from test data slipping into training, which makes results look better than they are.

Calibration: Whether predicted probabilities match reality: of borrowers scored 80%, do about 80% repay?

Feature contributions (SHAP): How much each factor pushed one borrower's prediction up or down.

Reject inference: The data only shows approved loans, so nothing is known about declined applicants.

PCA: Principal component analysis: combines many related variables into a few summary scores.

TF-IDF: Turns text into numbers by how important each word is in a document.

References

[1] Board of Governors of the Federal Reserve System. (2007, August). *Report to the Congress on credit scoring and its effects on the availability and affordability of credit*. https://www.federalreserve.gov/boarddocs/RptCongress/creditscore/general.htm

[2] Consumer Financial Protection Bureau. (n.d.). Regulation B (Equal Credit Opportunity Act), 12 CFR §1002.9: Notifications. https://www.consumerfinance.gov/rules-policy/regulations/1002/9/

[3] Basel Committee on Banking Supervision. (2004, June). Basel II: International convergence of capital measurement and capital standards: A revised framework. Bank for International Settlements. https://www.bis.org/publ/bcbs107.htm

[4] Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD Conference*. https://arxiv.org/abs/1603.02754

[5] Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems 30*. https://proceedings.neurips.cc/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf

[6] Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems 30*. https://arxiv.org/abs/1705.07874

[7] FinRegLab. (2021, September 16). The use of machine learning for credit underwriting: Market & data science context. finreglab.org (research report PDF)

[8] Consumer Financial Protection Bureau. (2022). Consumer Financial Protection Circular 2022-03: Adverse action notification requirements in connection with credit decisions based on complex algorithms. consumerfinance.gov. Withdrawn May 12, 2025: Morgan Lewis, CFPB guidance tracker: https://www.morganlewis.com/topics/cfpb-guidance-tracker-rescinded-and-remaining

[9] Institute of International Finance & EY. (2022, December). *IIF and EY survey report on machine learning: Uses in credit risk and AML applications* (public summary). https://www.iif.com/Publications/ID/5184/IIF-and-EY-2022-Survey-Report-on-Machine-Learning-Summary

[10] Kaggle. (2025). Predicting Loan Payback: Playground Series, Season 5, Episode 11. https://www.kaggle.com/competitions/playground-series-s5e11

[11] Kaufman, R. (2018, August 21). The history of the FICO® Score. myFICO. https://www.myfico.com/credit-education/blog/history-of-the-fico-score

[12] U.S. Department of the Treasury. (2024, December). *Artificial intelligence in financial services*. https://home.treasury.gov/system/files/136/Artificial-Intelligence-in-Financial-Services.pdf